

Programação Paralela com OpenMP

ELC139 - Programação Paralela

João Vicente Ferreira Lima (UFSM)

Universidade Federal de Santa Maria

`jvlima@inf.ufsm.br`

`http://www.inf.ufsm.br/~jvlima`

2023/1

Outline

- 1 Task parallelism
- 2 Data flow dependency

Outline

1 Task parallelism

- Task parallelism
- OpenMP

2 Data flow dependency

Task parallelism

Task parallelism

- **Task parallelism** or **functional parallelism** or **control parallelism**.
- Decomposes the **computation** rather than the **manipulated data**.
 - Programming model for tasks that perform different computations.
- Ex.: Cilk, Intel TBB, OpenMP.

Task dependency

- Tasks with dependencies can unfold a **directed acyclic graph** (DAG).
 - Expressed by synchronization such as `sync` keyword.
- If data dependencies are considered, the algorithm unfolds a **data flow graph** (DFG).
- Ex.: Jade, Athapascan, OpenMP (**new**), KAAPI/XKaapi, StarPU, OmpSs, Intel Offload.

Task parallelism

Task parallelism

- **Task parallelism** or **functional parallelism** or **control parallelism**.
- Decomposes the **computation** rather than the **manipulated data**.
 - Programming model for tasks that perform different computations.
- Ex.: Cilk, Intel TBB, OpenMP.

Task dependency

- Tasks with dependencies can unfold a **directed acyclic graph** (DAG).
 - Expressed by synchronization such as `sync` keyword.
- If data dependencies are considered, the algorithm unfolds a **data flow graph** (DFG).
- Ex.: Jade, Athapascan, OpenMP (**new**), KAAPI/XKaapi, StarPU, OmpSs, Intel Offload.

Task construct

```
#pragma omp task
```

Task construct

```
#pragma omp task
```

Barrier taskwait

```
#pragma omp taskwait
```

- Independent units of work.
- Recursive tasks.
- Unfold parallelism at runtime.
- The OpenMP implementation decides when/where to execute:
 - Immediately (in depth, **depth-first** or **work-first**)
 - Latter (in breadth, **breadth-first** or **help-first**)

OpenMP tasks

Task construct

```
#pragma omp task
```

Barrier taskwait

```
#pragma omp taskwait
```

- Independent units of work.
- Recursive tasks.
- Unfold parallelism at runtime.
- The OpenMP implementation decides when/where to execute:
 - Immediately (in depth, **depth-first** or **work-first**)
 - Latter (in breadth, **breadth-first** or **help-first**)

Linked list

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```

Linked list - parallel version

```
#pragma omp parallel
{
    #pragma omp single
    {
        node* p = head;
        while(p) {
            #pragma omp task firstprivate(p)
            process(p);
            p = p->next;
        }
    }
}
```

Cálculo de Fibonacci

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    x = fib( n - 1 );  
    y = fib( n - 2 );  
    return x + y;  
}
```

Cálculo de Fibonacci com OpenMP

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    #pragma omp task  
    x = fib( n - 1 );  
    #pragma omp task  
    y = fib( n - 2 );  
    #pragma omp taskwait  
    return x + y;  
}
```

Correto?

Não pois x e y são privados fora do escopo das tarefas.

Cálculo de Fibonacci com OpenMP

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    #pragma omp task  
    x = fib( n - 1 );  
    #pragma omp task  
    y = fib( n - 2 );  
    #pragma omp taskwait  
    return x + y;  
}
```

Correto?

Não pois x e y são privados fora do escopo das tarefas.

Cálculo de Fibonacci com OpenMP

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    #pragma omp task  
    x = fib( n - 1 );  
    #pragma omp task  
    y = fib( n - 2 );  
    #pragma omp taskwait  
    return x + y;  
}
```

Correto?

Não pois x e y são privados fora do escopo das tarefas.

Cálculo de Fibonacci com OpenMP

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    #pragma omp task shared(x)  
    x = fib( n - 1 );  
    #pragma omp task shared(y)  
    y = fib( n - 2 );  
    #pragma omp taskwait  
    return x + y;  
}
```

Agora sim

Necessitamos dos dois valores no cálculo.

2 Data flow dependency

- Data flow dependency
- Data flow example
- Data flow graph
- OpenMP data-flow tasks


Data flow dependency

Concept


Similar to a DAG of tasks, but combines **task dependencies** with **data-driven execution**.

- Execution is controlled by the **Data Flow Graph (DFG)**.
- Unlike the program recursion structure of **full strict model**.

Fully strict mode (Cilk)



Data flow graph



Universidade Federal de Santa Maria



Data access modes

- **Read only (RO or R)** - only read, no permission to modify.
- **Write only (WO or W)** - only write, no wait for data inputs.
- **Read and Write (RW)** - or exclusive mode, read and write.
- **Cumulative Write (CW)** - concurrent write and cumulative.

Data access modes

- **Read only (RO or R)** - only read, no permission to modify.
- **Write only (WO or W)** - only write, no wait for data inputs.
- **Read and Write (RW)** - or exclusive mode, read and write.
- **Cumulative Write (CW)** - concurrent write and cumulative.

Data access modes

- Read only (**RO** or **R**) - only read, no permission to modify.
- Write only (**WO** or **W**) - only write, no wait for data inputs.
- Read and Write (**RW**) - or exclusive mode, read and write.
- Cumulative Write (**CW**) - concurrent write and cumulative.

Data access modes

- Read only (**RO** or **R**) - only read, no permission to modify.
- Write only (**WO** or **W**) - only write, no wait for data inputs.
- Read and Write (**RW**) - or exclusive mode, read and write.
- Cumulative Write (**CW**) - concurrent write and cumulative.

Data flow dependency


The concepts here **are the same** from **computer architecture** in which we can replace **instruction** by **task**.

Data dependencies

- Task t_i produces a result that may be used by a task t_j .
- Task t_j is data dependent on task t_k , and t_k is data dependent on t_i ($t_i \rightarrow t_k \rightarrow t_j$).

```
compute( input, &result );
```

```
display( &result );
```





Data flow dependency

The concepts here **are the same** from **computer architecture** in which we can replace **instruction** by **task**.

Data dependencies

- Task t_i produces a result that may be used by a task t_j .
- Task t_j is data dependent on task t_k , and t_k is data dependent on t_i ($t_i \rightarrow t_k \rightarrow t_j$).

```
compute( input, &result );  
display( &result );
```




Data hazards

There is a **dependence** between task and the we must preserve the execution order.


- **RAW** (*Read after Write*).
- **WAW** (*Write after Write*).
- **WAR** (*Write after Read*).

Data flow dependency




- **RAW** (*Read after Write*) - or **true dependency**, **data dependency**.
- A task depends on the result produced by a previous task.

Data flow dependency



- **WAW** (*Write after Write*) - or **output dependency**.
- The execution order will affect the final output.

Data flow dependency



- **WAR** (*Write after Read*) - or **anti-dependency**.
- A task writes a value before it is read.

Data flow example

In our next examples, we will use the following keywords:


- **in** - read access.
- **out** - write access.
- **inout** - read and write access.

Data flow example

```
void reading(in int a) {}
void modifying(inout int b) {}
main(void)
{
    int a;
    reading( a );
    reading( a );
    reading( a );
    modifying( a );
    modifying( a );
}
```


Data flow example

```
void reading(in int a) {}  
void modifying(inout int b) {}  
main(void)  
{  
    int a;  
    reading( a );  
    reading( a );  
    reading( a );  
    modifying( a );  
    modifying( a );  
}
```




Data flow example

```
void reading(in int a) {}  
void modifying(inout int b) {}  
main(void)  
{  
    int a;  
    reading( a );  
    reading( a );  
    reading( a );  
    modifying( a );  
    modifying( a );  
}
```




Data flow example

```
void reading(in int a) {}  
void modifying(inout int b) {}  
main(void)  
{  
    int a;  
    reading( a );  
    reading( a );  
    reading( a );  
    modifying( a );  
    modifying( a );  
}
```




Data flow example

```
void reading(in int a) {}  
void modifying(inout int b) {}  
main(void)  
{  
    int a;  
    reading( a );  
    reading( a );  
    reading( a );  
    modifying( a );  
    modifying( a );  
}
```



Data flow example

```
void reading(in int a) {}  
void modifying(inout int b) {}  
main(void)  
{  
    int a;  
    reading( a );  
    reading( a );  
    reading( a );  
    modifying( a );  
    modifying( a );  
}
```



Fibonacci example

```
void fibo( int n, int* res )
{
    int x, y;
    if( n < 2 ){
        *res = n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
        *res = x+y;
    }
}

int main(void) {
    int n = 3, res;
    fibo( n, &res );
    print( res );
    return 0;
}
```

Notice

This recursive Fibonacci is not the best implementation, but it serves our purposes.

Dependency example (again)

In our next example, we will use the following keywords:

- **in** - read access.
- **out** - write access.
- **inout** - read and write access.
- **cout** - cumulative write with global reduction.

Notice

This recursive Fibonacci is not the best implementation, but it serves our purposes.

Dependency example (again)

In our next example, we will use the following keywords:

- **in** - read access.
- **out** - write access.
- **inout** - read and write access.
- **cout** - cumulative write with global reduction.

Fibonacci example

```
void fibo( in int n, out int* res )
{
    int x, y;
    if( n < 2 ){
        *res = n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
        *res = x+y;
    }
}

int main(void) {
    int n = 3, res;
    fibo( n, &res );
    print( res );
    return 0;
}
```

Fibonacci example

```
void fibo( in int n, out int* res )
{
    int x, y;
    if( n < 2 ){
        *res = n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
        *res = x+y;
    }
}

int main(void) {
    int n = 3, res;
    fibo( n, &res );
    print( res );
    return 0;
}
```

Fibonacci example

Previous Fibonacci example

```
} else {  
    fibo( n-1, &x );  
    fibo( n-2, &y );  
    *res = x+y;  
}
```

Synchronization problem

If our tasks execute in parallel, we would like **to wait** for the results from the previous two `fibo` tasks.

Solution

- 1 An explicit synchronization (Cilk's style).
- 2 A task that depends on the results from the two `fibo` tasks.

Fibonacci example

Previous Fibonacci example

```
} else {  
    fibo( n-1, &x );  
    fibo( n-2, &y );  
    *res = x+y;  
}
```

Synchronization problem

If our tasks execute in parallel, we would like **to wait** for the results from the previous two `fibo` tasks.

Solution

- 1 An explicit synchronization (Cilk's style).
- 2 A task that depends on the results from the two `fibo` tasks.

Fibonacci example

Previous Fibonacci example

```
} else {  
    fibo( n-1, &x );  
    fibo( n-2, &y );  
    *res = x+y;  
}
```

Synchronization problem

If our tasks execute in parallel, we would like **to wait** for the results from the previous two `fibo` tasks.

Solution

- 1 An explicit synchronization (Cilk's style).
- 2 A task that depends on the results from the two `fibo` tasks.

Fibonacci example

```
void sum( out int* res, in int x, in int y )
{
    *a = x + y;
}

void fibo( in int n, out int* res )
{
    int x, y;
    if( n < 2 ){
        *res = n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
        sum( res, x, y );
    }
}
```

Fibonacci example

```
void sum( out int* res, in int x, in int y )
{
    *a = x + y;
}

void fibo( in int n, out int* res )
{
    int x, y;
    if( n < 2 ){
        *res = n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
        sum( res, x, y );
    }
}
```

Fibonacci example (cumulative)

```
void fibo( in int n, cout int* res )
{
    int x, y;
    if( n < 2 ){
        *res += n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
    }
}
```


Fibonacci example (cumulative)

```
void fibo( in int n, cout int* res )
{
    int x, y;
    if( n < 2 ){
        *res += n;
    } else {
        fibo( n-1, &x );
        fibo( n-2, &y );
    }
}
```


Data flow graph

- Data flow graph (DFG) combines task dependencies with data driven execution.


Fully strict mode (Cilk)




Data flow graph




DAG of Fibonacci $n = 3$



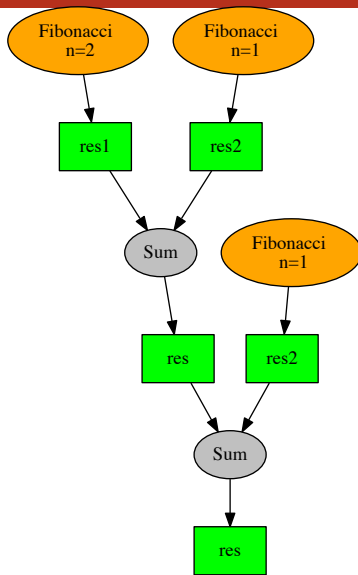
DAG of Fibonacci $n = 3$



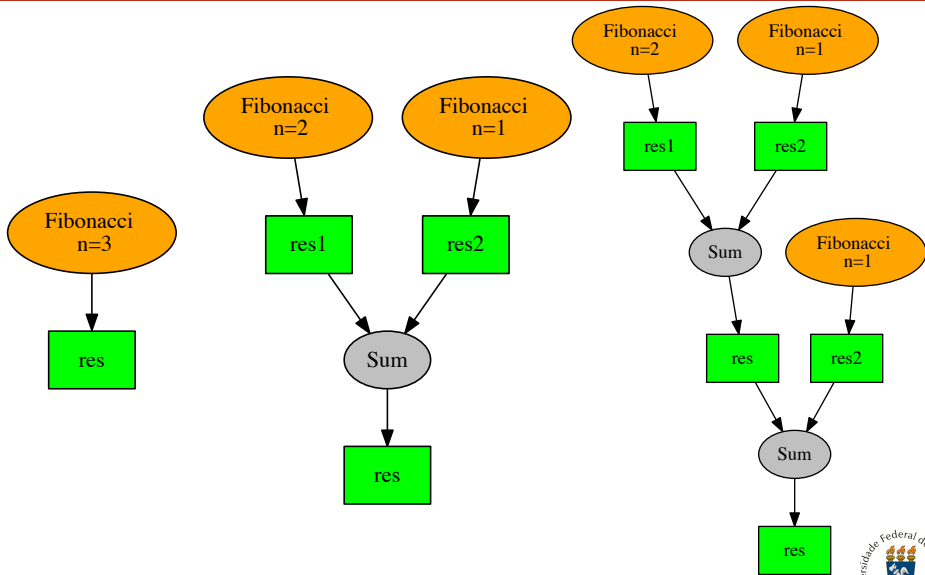
DFG of Fibonacci $n = 3$



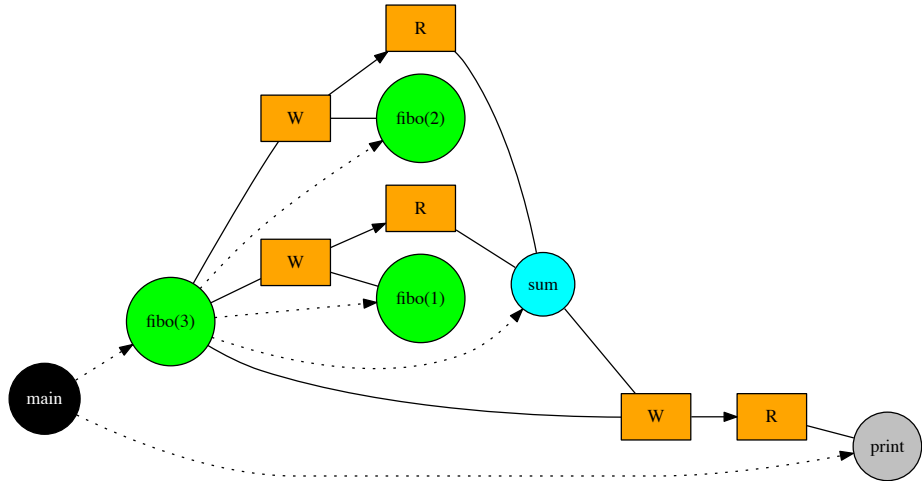
DFG of Fibonacci $n = 3$




DFG of Fibonacci $n = 3$




DFG of Fibonacci $n = 3$



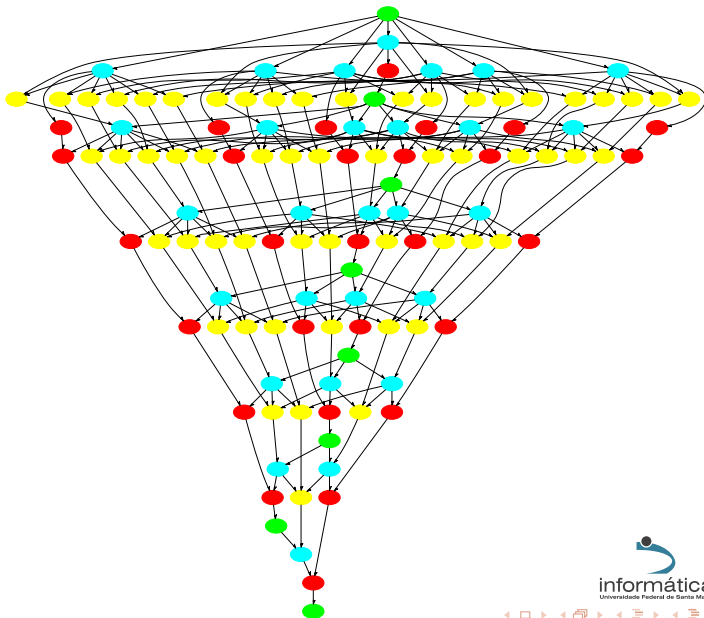
DFG of Fibonacci $n = 3$ (cumulative)



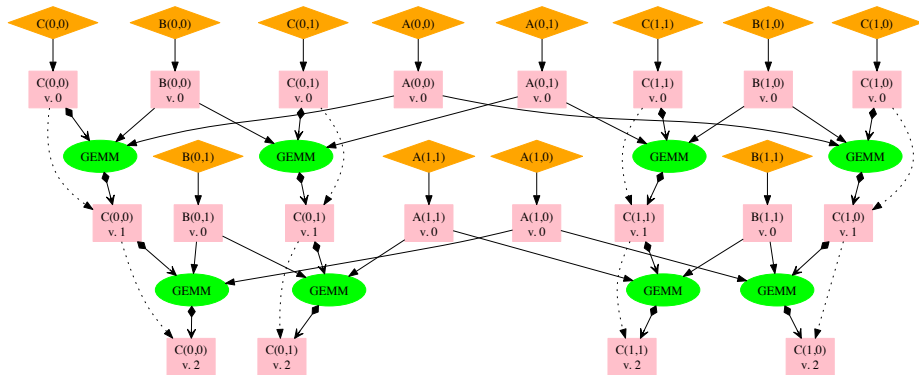
DFG of Cholesky factorization



DFG of Cholesky factorization



DFG of Blocked matrix multiplication



Dependências de Dados

- OpenMP 4.0 inclui dependências de dados entre tarefas (`task`)
- Diretiva `depend`
 - `in` – dados de entrada
 - `out` – dados de saída
 - `inout` – dados de entrada e saída
 - `depobj` – dependências manuais (`omp_depend_t`)
- Sincronização recursiva pela construção `taskgroup`
 - Sincroniza um bloco de código
 - `taskwait` **sincroniza tarefas criadas no mesmo nível**

```
#pragma omp taskgroup
{
    #pragma omp task depend(in:dados) depend(out:saida)
    foo(dados, saida);
}
```

Dependências de Dados


- OpenMP 4.0 inclui dependências de dados entre tarefas (task)
- Diretiva `depend`
 - `in` – dados de entrada
 - `out` – dados de saída
 - `inout` – dados de entrada e saída
 - `depobj` – dependências manuais (`omp_depend_t`)
- Sincronização recursiva pela construção `taskgroup`
 - Sincroniza um bloco de código
 - `taskwait` sincroniza tarefas criadas no mesmo nível

```
#pragma omp taskgroup
```

```
{
  #pragma omp task depend(in:dados) depend(out:saida)
  foo(dados, saida);
}
```

Dependências de Dados

```
1 int x = 0;
2 #pragma omp parallel
3 #pragma omp single
4 {
5     #pragma omp task depend(inout: x) // T1
6     { .... }
7
8     #pragma omp task depend(in: x) // T2
9     { .... }
10
11    #pragma omp task depend(in: x) // T3
12    { .... }
13
14    #pragma omp task depend(inout: x) // T4
15    { .... }
16 }
```



Fibonacci com OpenMP depend

```
int fib( int n ) {  
    int x, y;  
    if( n < 2 ) return n;  
    #pragma omp taskgroup  
    {  
        #pragma omp task shared(x) depend(in:n) depend(out:x)  
        x = fib( n - 1 );  
        #pragma omp task shared(y) depend(in:n) depend(out:y)  
        y = fib( n - 2 );  
    }  
    return x + y;  
}
```

Espera por dados com `taskwait`

```
int x = 0, y = 0;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task depend(inout: x)
    x++;

    #pragma omp task depend(in: y)
    std::cout << y << std::endl;

    #pragma omp taskwait depend(in: x)
    std::cout << x << std::endl;
}
```

Dependências de Dados

- Dependências manuais com `depobj`
- Permite dependências complexas.
- Novo tipo opaco `omp_depend_t`

```
int x = 0, y = 0;
#pragma omp parallel
#pragma omp single
{
    omp_depend_t obj;
    #pragma omp depobj(obj) depend(inout: x)

    #pragma omp task depend(depobj: obj) // T1
    x++;

    #pragma omp depobj(obj) update(in)

    #pragma omp task depend(depobj: obj) // T2
    std::cout << y << std::endl;

    #pragma omp depobj(obj) destroy
}
```

Multiplicação de matrizes $C = A \times B + C$

```
void gemm_omp(double *A, double *B, double *C, int n)
{
    #pragma omp parallel
    {
        int i, j, k;
        #pragma omp for
        for (i = 0; i < n; i++) {
            for (j = 0; j < n; j++) {
                for (k = 0; k < n; k++) {
                    C[i*n+j] += A[i*n+k]*B[k*n+j];
                }
            }
        }
    }
}
```

Multiplicação de matrizes $C = Ax + B$

```
void matmul(int NB, float A[NB][NB], float B[NB][NB], float  
    ↪ C[NB][NB])  
{  
    #pragma omp parallel  
    #pragma omp single  
    {  
        for(int i = 0; i < NB; i++)  
            for(int j = 0; j < NB; j++)  
                for(int k = 0; k < NB; k++)  
#pragma omp task depend(in:A[i][k],B[k][j])  
    ↪    depend(inout:C[i][j])  
                matmul_tile( A[i][k], B[k][j], C[i][j] );  
    }  
}
```

<https://joao-ufsm.github.io/par2023a/>

